

Docker Infrastructure Documentation

Overview

This document describes the enhanced Docker infrastructure setup for the PowerShell Script Analysis Platform, including connection pooling, high availability, and automated backups.

Architecture Components

1. PgBouncer Connection Pooling

Purpose: Manages PostgreSQL connections efficiently to prevent connection exhaustion and improve performance.

Configuration: - Port: 6432 - Pool Mode: Transaction - Max Client Connections: 1000 - Default Pool Size: 25 connections per database - Reserve Pool: 5 connections

Configuration Files: - `/docker/pgbouncer/pgbouncer.ini` - Main configuration - `/docker/pgbouncer/userlist.txt` - User authentication

Benefits: - Reduces database connection overhead - Prevents connection exhaustion - Improves application scalability - Faster connection establishment

Usage:

```
# Connect via pgBouncer instead of direct PostgreSQL
DB_HOST=pgbouncer
DB_PORT=6432
```

2. Redis Cluster with Sentinel

Purpose: Provides high availability for caching and session storage with automatic failover.

Architecture: - 1 Redis Master (port 6379) - 2 Redis Replicas (ports 6380, 6381) - 3 Redis Sentinels (ports 26379, 26380, 26381)

Features: - Automatic failover detection - Master election on failure - Data replication across replicas - Quorum-based decision making (2/3 sentinels)

Configuration Files: - `/docker/redis/redis-master.conf` - Master configuration - `/docker/redis/redis-replica.conf` - Replica configuration - `/docker/redis/sentinel.conf` - Sentinel monitoring configuration

Failover Process: 1. Sentinel detects master failure (5 second timeout) 2. Quorum agreement among sentinels (2 out of 3) 3. Automatic promotion of replica to master 4. Client reconnection to new master

3. Backup Automation Service

Purpose: Automated backup, retention, and disaster recovery for PostgreSQL and Redis.

Features: - Scheduled full and incremental backups - Automated cleanup based on retention policy - Optional S3 cloud storage integration - Health monitoring and alerting - Backup verification and metadata tracking

Backup Schedules (configurable via environment variables): - PostgreSQL Full Backup: Daily at 2 AM - PostgreSQL Incremental: Every 6 hours - Redis Snapshot: Every 4 hours - Cleanup Old Backups: Daily at 3 AM - Health Checks: Every 5 minutes

Backup Scripts: - `postgres-backup.sh` - PostgreSQL backup (full/incremental) - `redis-backup.sh` - Redis snapshot backup - `cleanup-backups.sh` - Remove old backups based on retention policy - `health-check.sh` - Monitor service health - `restore-postgres.sh` - Restore PostgreSQL from backup - `restore-redis.sh` - Restore Redis from backup

Backup Location: - Local: `/backups/postgres` and `/backups/redis` - Cloud (optional): S3 bucket (configurable)

Services Overview

Service	Port(s)	Purpose	Health Check
frontend	3000	React UI	HTTP GET /
backend	4000	API Server	HTTP GET /health
ai-service	8000	AI Analysis	TCP
postgres	5432	Database	pg_isready
pgbouncer	6432	Connection Pool	pg_isready
redis-master	6379	Cache Master	PING
redis-replica-1	6380	Cache Replica	PING
redis-replica-2	6381	Cache Replica	PING
redis-sentinel-1	26379	Failover Monitor	TCP
redis-sentinel-2	26380	Failover Monitor	TCP
redis-sentinel-3	26381	Failover Monitor	TCP
backup-service	-	Backup Automation	Internal

Network Configuration

Network: psscript-network - Type: Bridge - Subnet: 172.25.0.0/16 - DNS:
Automatic service discovery by name

Volume Management

Persistent Volumes: - `postgres_data` - PostgreSQL data directory -
`redis_master_data` - Redis master data - `redis_replica_1_data` - Redis
replica 1 data - `redis_replica_2_data` - Redis replica 2 data -
`redis_sentinel_1_data` - Sentinel 1 configuration - `redis_sentinel_2_data`
- Sentinel 2 configuration - `redis_sentinel_3_data` - Sentinel 3 configuration

Bind Mounts: - `./backups` - Backup storage directory - `./docker/pgbouncer` -
PgBouncer configuration - `./docker/redis` - Redis configuration -
`./docker/postgres` - PostgreSQL configuration

Getting Started

Initial Setup

1. Copy environment file:

```
cp .env.example .env  
# Edit .env with your configuration
```

1. Create required directories:

```
mkdir -p backups/postgres backups/redis backups/logs
```

1. Start all services:

```
docker-compose up -d
```

1. Verify services are running:

```
docker-compose ps
```

Service Management

Start all services:

```
docker-compose up -d
```

Stop all services:

```
docker-compose down
```

View logs:

```
# All services  
docker-compose logs -f  
  
# Specific service  
docker-compose logs -f backend  
docker-compose logs -f redis-master  
docker-compose logs -f pgbouncer
```

Restart a service:

```
docker-compose restart backend
```

Scale replica count (if needed):

```
docker-compose up -d --scale redis-replica=3
```

Monitoring

Check PgBouncer status:

```
docker-compose exec pgbouncer psql -h localhost -p 6432 -U postgres  
docker-compose exec pgbouncer psql -h localhost -p 6432 -U postgres
```

Check Redis Sentinel status:

```
docker-compose exec redis-sentinel-1 redis-cli -p 26379 SENTINEL  
docker-compose exec redis-sentinel-1 redis-cli -p 26379 SENTINEL
```

Check Redis replication:

```
docker-compose exec redis-master redis-cli INFO replication  
docker-compose exec redis-replica-1 redis-cli INFO replication
```

View backup status:

```
# Check backup logs
tail -f backups/logs/backup.log

# View backup inventory
ls -lh backups/postgres/
ls -lh backups/redis/

# Check health status
tail -f backups/logs/health.log
```

Backup and Recovery

Manual Backups

PostgreSQL Full Backup:

```
docker-compose exec backup-service /scripts/postgres-backup.sh full
```

PostgreSQL Incremental Backup:

```
docker-compose exec backup-service /scripts/postgres-backup.sh incremental
```

Redis Backup:

```
docker-compose exec backup-service /scripts/redis-backup.sh
```

Restore Operations

Restore PostgreSQL:

```
# List available backups
docker-compose exec backup-service ls -lh /backups/postgres/

# Restore from backup
docker-compose exec backup-service /scripts/restore-postgres.sh /backups/postgres/2023-10-10T14:00:00Z.bak
```

Restore Redis:

```
# List available backups
docker-compose exec backup-service ls -lh /backups/redis/

# Restore from backup
docker-compose exec backup-service /scripts/restore-redis.sh /backups/redis/2023-10-10T14:00:00Z.bak
```

S3 Cloud Backup Configuration

Add to your `.env` file:

```
BACKUP_S3_BUCKET=your-bucket-name  
BACKUP_S3_REGION=us-east-1  
AWS_ACCESS_KEY_ID=your-access-key  
AWS_SECRET_ACCESS_KEY=your-secret-key
```

Backups will automatically upload to S3 after local backup completes.

Disaster Recovery

Complete System Recovery

1. Stop all services:

```
docker-compose down
```

1. Restore PostgreSQL:

```
docker-compose up -d postgres pgbouncer backup-service
docker-compose exec backup-service /scripts/restore-postgres.sh /backups/postgres/
```

1. Restore Redis:

```
docker-compose up -d redis-master
docker-compose exec backup-service /scripts/restore-redis.sh /backups/redis/
```

1. Start remaining services:

```
docker-compose up -d
```

Testing Failover

Test Redis Failover:

```
# Stop master
docker-compose stop redis-master

# Watch sentinel logs
docker-compose logs -f redis-sentinel-1

# Verify new master
docker-compose exec redis-sentinel-1 redis-cli -p 26379 SENTINEL m

# Restart original master (becomes replica)
docker-compose start redis-master
```

Performance Tuning

PgBouncer Optimization

Edit `/docker/pgbouncer/pgbouncer.ini`:

```
# Increase pool sizes for high-traffic applications
default_pool_size = 50
max_client_conn = 2000

# Adjust timeouts
query_wait_timeout = 60
```

PostgreSQL Optimization

Edit `/docker/postgres/postgresql.conf`:

```
# Increase memory for better performance
shared_buffers = 512MB
effective_cache_size = 2GB
work_mem = 32MB
```

Redis Optimization

Edit `/docker/redis/redis-master.conf`:

```
# Increase memory limit
maxmemory 1gb

# Adjust eviction policy
maxmemory-policy allkeys-lfu
```

Troubleshooting

PgBouncer Connection Issues

```
# Check PgBouncer logs
docker-compose logs pgbouncer

# Test direct connection
docker-compose exec pgbouncer psql -h postgres -p 5432 -U postgres

# Check pool status
docker-compose exec pgbouncer psql -h localhost -p 6432 -U postgres
```

Redis Failover Not Working

```
# Check sentinel logs
docker-compose logs redis-sentinel-1

# Verify sentinel configuration
docker-compose exec redis-sentinel-1 cat /usr/local/etc/redis/sentinel.conf

# Check master status
docker-compose exec redis-sentinel-1 redis-cli -p 26379 SENTINEL m
```

Backup Failures

```
# Check backup service logs
docker-compose logs backup-service
tail -f backups/logs/backup.log

# Verify disk space
df -h

# Test manual backup
docker-compose exec backup-service /scripts/postgres-backup.sh full
```

Security Considerations

1. **Change default passwords** in production:
2. PostgreSQL password
3. Redis password (uncomment in redis configs)
4. PgBouncer user password
5. **Enable SSL/TLS:**
6. Configure PostgreSQL SSL
7. Configure Redis TLS
8. Use encrypted connections in pgBouncer
9. **Restrict network access:**
10. Use firewall rules
11. Limit exposed ports
12. Use VPN for remote access
13. **Encrypt backups:**
14. Enable S3 encryption
15. Use encrypted volumes
16. Secure backup transfer channels

Maintenance

Regular Maintenance Tasks

1. Monitor disk usage:

```
docker-compose exec backup-service df -h /backups
```

1. Check service health:

```
docker-compose ps
docker-compose exec backup-service /scripts/health-check.sh
```

1. Review logs:

```
docker-compose logs --tail=100 postgres
docker-compose logs --tail=100 redis-master
```

1. Update Docker images:

```
docker-compose pull
docker-compose up -d
```

1. Vacuum PostgreSQL:

```
docker-compose exec postgres vacuumdb -U postgres -d psscript --a
```

Support and Resources

- Docker Compose Documentation: <https://docs.docker.com/compose/>
- PgBouncer Documentation: <https://www.pgbouncer.org/>
- Redis Sentinel Documentation: <https://redis.io/topics/sentinel>
- PostgreSQL Documentation: <https://www.postgresql.org/docs/>

Generated 2026-01-13 06:26 UTC