

LangGraph 1.0 Production Implementation

Overview

This document provides comprehensive information about the LangGraph 1.0 production orchestrator implementation for PowerShell script analysis.

Architecture

Components



State Schema

```
class PowerShellAnalysisState(TypedDict):
    # Message history with automatic deduplication
    messages: Annotated[Sequence[BaseMessage], add_messages]

    # Script content being analyzed
    script_content: Optional[str]

    # Analysis results from various stages
    analysis_results: Dict[str, Any]

    # Current workflow stage
    current_stage: str

    # Security findings
    security_findings: List[Dict[str, Any]]

    # Code quality metrics
    quality_metrics: Dict[str, float]

    # Optimization recommendations
    optimizations: List[Dict[str, Any]]

    # Error tracking
    errors: List[Dict[str, Any]]

    # Metadata
    workflow_id: str
    started_at: str
    completed_at: Optional[str]

    # Human-in-the-loop flags
    requires_human_review: bool
    human_feedback: Optional[str]

    # Final output
    final_response: Optional[str]
```

API Reference

Analyze Script

Analyze a PowerShell script using the LangGraph orchestrator.

Endpoint: POST /langgraph/analyze

Request Body:

```
{  
    "script_content": "Get-Process | Where-Object CPU -gt 100",  
    "thread_id": "optional_thread_id",  
    "require_human_review": false,  
    "stream": false,  
    "model": "gpt-4",  
    "api_key": "optional_api_key"  
}
```

Response:

```
{  
    "workflow_id": "analysis_1704649200.123",  
    "status": "completed",  
    "final_response": "This script retrieves processes with CPU usage  
    > 5% and sorts them by memory usage.",  
    "analysis_results": {  
        "analyze_powershell_script": {  
            "purpose": "Process monitoring",  
            "complexity": "Low",  
            "line_count": 1  
        },  
        "security_scan": {  
            "risk_level": "LOW",  
            "risk_score": 0,  
            "findings": []  
        },  
        "quality_analysis": {  
            "quality_score": 6.0,  
            "metrics": {  
                "total_lines": 1,  
                "code_lines": 1  
            }  
        }  
    },  
    "current_stage": "completed",  
    "started_at": "2026-01-07T12:00:00.000Z",  
    "completed_at": "2026-01-07T12:00:04.523Z",  
    "requires_human_review": false  
}
```

Provide Human Feedback

Continue a paused workflow with human feedback.

Endpoint: POST /langgraph/feedback

Request Body:

```
{  
  "thread_id": "analysis_1704649200.123",  
  "feedback": "The security analysis looks good. Please proceed with caution."  
}
```

Response: Same format as analyze response

Health Check

Check the health of the LangGraph service.

Endpoint: GET /langgraph/health

Response:

```
{  
  "status": "healthy",  
  "service": "LangGraph Production Orchestrator",  
  "version": "1.0.5",  
  "features": {  
    "checkpointing": true,  
    "human_in_the_loop": true,  
    "streaming": true,  
    "durable_execution": true  
  },  
  "model": "gpt-4",  
  "checkpointer_type": "MemorySaver"  
}
```

Service Info

Get detailed information about the orchestrator.

Endpoint: GET /langgraph/info

Response:

```
{  
  "orchestrator": "LangGraph Production Orchestrator",  
  "version": "1.0.5",  
  "description": "Production-grade PowerShell script analysis using AI",  
  "workflow_stages": ["analyze", "tools", "synthesis", "human_review"],  
  "available_tools": [  
    {  
      "name": "analyze_powershell_script",  
      "description": "Analyze script purpose, functionality, and security",  
    },  
    {  
      "name": "security_scan",  
      "description": "Comprehensive security analysis and vulnerability detection",  
    },  
    {  
      "name": "quality_analysis",  
      "description": "Code quality evaluation and best practices compliance",  
    },  
    {  
      "name": "generate_optimizations",  
      "description": "Generate optimization recommendations for performance improvement",  
    }  
  "supported_models": ["gpt-4", "gpt-4-turbo", "gpt-3.5-turbo"]  
}
```

Batch Analysis

Analyze multiple scripts concurrently.

Endpoint: POST /langgraph/batch-analyze

Request Body:

```
{  
  "scripts": [  
    "Get-Process",  
    "Get-Service | Where-Object Status -eq 'Running'"  
  ]  
}
```

Response:

```
{  
  "total": 2,  
  "successful": 2,  
  "failed": 0,  
  "results": [  
    {  
      "index": 0,  
      "workflow_id": "analysis_1704649200.123",  
      "status": "completed"  
    },  
    {  
      "index": 1,  
      "workflow_id": "analysis_1704649200.456",  
      "status": "completed"  
    }  
  ],  
  "errors": []  
}
```

Test Orchestrator

Test the orchestrator with a sample script.

Endpoint: POST /langgraph/test

Response:

```
{  
  "test_status": "passed",  
  "result": {  
    "workflow_id": "...",  
    "status": "completed",  
    ...  
  }  
}
```

Tools

analyze_powershell_script

Analyzes script purpose, structure, and basic metrics.

Input: Script content string

Output:

```
{  
  "purpose": "File search and filtering",  
  "complexity": "Medium",  
  "parameters": {"Path": "string"},  
  "functions": [],  
  "line_count": 7,  
  "timestamp": "2026-01-07T12:00:00.000Z"  
}
```

security_scan

Performs comprehensive security analysis.

Input: Script content string

Output:

```
{  
    "risk_level": "MEDIUM",  
    "risk_score": 12,  
    "findings": [  
        {  
            "category": "Network Activity",  
            "severity": 5,  
            "pattern": "invoke-webrequest",  
            "description": "Makes web requests"  
        }  
    ],  
    "findings_count": 1,  
    "best_practices": ["Uses error handling", "Uses parameter validation"],  
    "timestamp": "2026-01-07T12:00:00.000Z"  
}
```

Security Patterns Detected: - invoke-expression - Code injection risk (severity: 10) - downloadstring - Remote code execution (severity: 9) - bypass - Security control bypass (severity: 8) - encodedcommand - Obfuscation (severity: 8) - hidden - Stealth execution (severity: 7) - downloadfile - Untrusted download (severity: 7) - start-process - Process creation (severity: 6) - add-type - Code compilation (severity: 6) - invoke-webrequest - Network activity (severity: 5)

Risk Levels: - CRITICAL: risk_score > 30 - HIGH: risk_score > 20 - MEDIUM: risk_score > 10 - LOW: risk_score ≤ 10

quality_analysis

Evaluates code quality and best practices.

Input: Script content string

Output:

```
{  
    "quality_score": 7.5,  
    "metrics": {  
        "total_lines": 50,  
        "comment_lines": 10,  
        "empty_lines": 5,  
        "code_lines": 35,  
        "comment_ratio": 0.286  
    },  
    "issues": ["Script is very long - consider breaking into modules"],  
    "recommendations": [  
        "Add comment-based help",  
        "Implement try/catch error handling"  
    ],  
    "timestamp": "2026-01-07T12:00:00.000Z"  
}
```

Quality Scoring: - Base score: 5.0 - +1.0 for CmdletBinding - +0.5 for param block
- +0.5 for comments (>10% ratio) - +1.0 for try/catch blocks - -0.5 for very long scripts (>500 lines) - -0.5 for many long lines (>5 lines over 120 chars)

generate_optimizations

Generates optimization recommendations.

Input: Script content and quality metrics

Output:

```
{  
  "total_optimizations": 3,  
  "optimizations": [  
    {  
      "category": "Performance",  
      "priority": "Medium",  
      "recommendation": "Use .ForEach() method instead of ForEach-",  
      "impact": "Can improve loop performance by 2-3x"  
    },  
    {  
      "category": "Reliability",  
      "priority": "High",  
      "recommendation": "Add try/catch blocks for error handling",  
      "impact": "Prevents script failures and improves debugging"  
    },  
    {  
      "category": "Documentation",  
      "priority": "Medium",  
      "recommendation": "Add comment-based help",  
      "impact": "Improves code understanding and maintenance"  
    }  
  "timestamp": "2026-01-07T12:00:00.000Z"  
}
```

Workflow

Standard Analysis Flow

1. **START** → Initial state created
2. **analyze** → LLM determines required analysis
3. **tools** → Execute requested tools (security_scan, quality_analysis, etc.)
4. **analyze** → LLM processes tool results
5. **synthesis** → Generate final comprehensive response
6. **END** → Workflow complete

Human-in-the-Loop Flow

1. **START** → Initial state created with `requires_human_review=true`
2. **analyze** → LLM determines required analysis
3. **tools** → Execute requested tools
4. **human_review** → Workflow pauses, waiting for feedback
5. *[Human provides feedback via /feedback endpoint]*
6. **analyze** → Re-analyze with human input
7. **synthesis** → Generate final response
8. **END** → Workflow complete

Checkpointing

Development (MemorySaver)

```
orchestrator = LangGraphProductionOrchestrator(  
    api_key="...",  
    use_postgres_checkpointing=False  
)
```

- Stores state in memory
- Fast for development/testing
- No persistence across restarts
- Suitable for development environments

Production (PostgresSaver)

```
orchestrator = LangGraphProductionOrchestrator(  
    api_key="...",  
    use_postgres_checkpointing=True,  
    postgres_connection_string="postgresql://user:pass@host/db"  
)
```

- Stores state in PostgreSQL
- Durable across restarts
- Supports state recovery
- Required for production

Database Schema:

```
CREATE TABLE checkpoints (
    thread_id TEXT PRIMARY KEY,
    checkpoint_id TEXT NOT NULL,
    parent_checkpoint_id TEXT,
    checkpoint JSONB NOT NULL,
    metadata JSONB,
    created_at TIMESTAMP DEFAULT NOW()
);

CREATE INDEX idx_checkpoints_thread ON checkpoints(thread_id);
CREATE INDEX idx_checkpoints_parent ON checkpoints(parent_checkpoint_id);
```

Configuration

Environment Variables

```
# OpenAI API Key (required)
OPENAI_API_KEY=sk-...

# Model selection
DEFAULT_MODEL=gpt-4

# Checkpointing (production)
USE_POSTGRES_CHECKPOINTING=true
DATABASE_URL=postgresql://user:pass@host:5432/psscript

# Logging
LOG_LEVEL=INFO

# Feature flags
ENABLE_LANGGRAPH=true
LANGGRAPH_TRAFFIC_PERCENTAGE=100
```

Code Configuration

```
# Initialize orchestrator
orchestrator = LangGraphProductionOrchestrator(
    api_key=os.getenv("OPENAI_API_KEY"),
    model="gpt-4",
    use_postgres_checkpointing=True,
    postgres_connection_string=os.getenv("DATABASE_URL")
)

# Analyze with options
result = await orchestrator.analyze_script(
    script_content=script,
    thread_id="my-session-id",
    require_human_review=False,
    stream=False
)
```

Monitoring

Metrics to Track

1. Performance Metrics

2. Workflow duration (p50, p95, p99)
3. Tool execution time
4. LLM response time
5. Total analysis time

6. Reliability Metrics

7. Success rate
8. Error rate by type
9. Retry count
10. Checkpoint recovery success

11. Resource Metrics

12. Memory usage
13. State size
14. Database query count
15. API calls per analysis

16. Business Metrics

17. Analyses per day
18. Average tools used per analysis
19. Human review request rate
20. User satisfaction

Logging

The orchestrator uses structured logging:

```
logger.info(f"Entering analyze_node for workflow {workflow_id}")  
logger.warning("Human review required - workflow paused")  
logger.error(f"Error in analyze_script: {e}", exc_info=True)
```

Log Levels: - DEBUG : Detailed state transitions - INFO : Workflow milestones -
WARNING : Recoverable issues - ERROR : Failures requiring attention

Error Handling

Automatic Recovery

The orchestrator handles errors gracefully:

```
try:  
    result = await orchestrator.analyze_script(script)  
except Exception as e:  
    # Error logged and returned in response  
    return {  
        "error": str(e),  
        "workflow_id": workflow_id,  
        "status": "failed"  
    }
```

Retry Logic

Tool execution includes automatic retries:

```
@tool  
def security_scan(script_content: str) -> str:  
    try:  
        # Analysis logic  
        return json.dumps(result)  
    except Exception as e:  
        logger.error(f"Error in security_scan: {e}")  
        return json.dumps({"error": str(e)})
```

State Recovery

With checkpointing enabled, workflows can resume after failures:

```
# Original execution fails
result = await orchestrator.analyze_script(
    script_content=script,
    thread_id="session-123"
)

# Resume from checkpoint after restart
orchestrator = LangGraphProductionOrchestrator()
result = await orchestrator.analyze_script(
    script_content=script,
    thread_id="session-123"  # Same thread ID
)
# Continues from last checkpoint
```

Best Practices

1. Always Use Thread IDs for Sessions

```
# Good
result = await orchestrator.analyze_script(
    script_content=script,
    thread_id=user_session_id
)

# Bad - generates random ID
result = await orchestrator.analyze_script(
    script_content=script
)
```

2. Enable Checkpointing in Production

```
# Production
orchestrator = LangGraphProductionOrchestrator(
    use_postgres_checkpointing=True,
    postgres_connection_string=DATABASE_URL
)

# Development only
orchestrator = LangGraphProductionOrchestrator()
```

3. Handle Streaming Properly

```
if stream:
    async for event in orchestrator.graph.astream(state, config):
        # Process each event
        yield event
```

4. Use Human Review for Sensitive Scripts

```
# For scripts with high risk scores
result = await orchestrator.analyze_script(
    script_content=script,
    require_human_review=True if risk_score > 50 else False
)
```

5. Monitor and Alert

```
# Track metrics
duration = result["completed_at"] - result["started_at"]
if duration > 10: # seconds
    alert("Slow analysis detected")

# Track errors
if result["status"] == "failed":
    alert(f"Analysis failed: {result['error']}")
```

Troubleshooting

Issue: Workflow hangs in tool execution

Symptoms: Analysis doesn't complete, stuck in "tool_execution" stage

Diagnosis:

```
# Check state
logger.info(f"Current stage: {state['current_stage']}")
logger.info(f"Last message: {state['messages'][-1]}")
```

Solution: Increase timeout or check tool implementation

Issue: Checkpoint not found

Symptoms: KeyError: 'checkpoint_id'

Diagnosis: Thread ID doesn't exist or checkpointing not enabled

Solution:

```
# Ensure checkpointing is enabled
orchestrator = LangGraphProductionOrchestrator(
    use_postgres_checkpointing=True,
    postgres_connection_string=DATABASE_URL
)
```

Issue: High memory usage

Symptoms: Memory grows during execution

Diagnosis: State accumulation in long-running workflows

Solution: Limit message history or use PostgreSQL checkpointing

```
# Trim old messages from state
state["messages"] = state["messages"][-10:] # Keep last 10
```

Issue: Tool execution errors

Symptoms: Tools return error responses

Diagnosis: Check tool logs and input validation

Solution:

```
# Add input validation
@tool
def analyze_powershell_script(script_content: str) -> str:
    if not script_content or len(script_content) < 10:
        return json.dumps({"error": "Invalid script content"})
    # Continue with analysis
```

Performance Optimization

1. Parallel Tool Execution

Tools are executed sequentially by default. For independent tools:

```
# Future enhancement: parallel tool execution
results = await asyncio.gather(
    security_scan(script),
    quality_analysis(script)
)
```

2. Caching

Cache frequent analysis results:

```
from functools import lru_cache

@lru_cache(maxsize=1000)
def get_cached_analysis(script_hash: str) -> Dict:
    # Return cached result if available
    pass
```

3. Model Selection

Use faster models for simple scripts:

```
# Simple script - use faster model
if len(script.split('\n')) < 20:
    model = "gpt-3.5-turbo"
else:
    model = "gpt-4"
```

4. Batch Processing

Process multiple scripts together:

```
results = await asyncio.gather(*[
    orchestrator.analyze_script(script)
    for script in scripts
])
```

Testing

Unit Tests

```
import pytest
from agents.langgraph_production import (
    analyze_powershell_script,
    security_scan,
    quality_analysis
)

def test_security_scan_dangerous_script():
    script = "Invoke-Expression $userInput"
    result = json.loads(security_scan(script))
    assert result["risk_level"] == "CRITICAL"
    assert len(result["findings"]) > 0

def test_quality_analysis():
    script = "[CmdletBinding()]\nparam($Path)\nGet-Item $Path"
    result = json.loads(quality_analysis(script))
    assert result["quality_score"] >= 5.0
```

Integration Tests

```
@pytest.mark.asyncio
async def test_full_workflow():
    orchestrator = LangGraphProductionOrchestrator()
    script = "Get-Process | Select-Object Name, CPU"

    result = await orchestrator.analyze_script(script)

    assert result["status"] == "completed"
    assert result["final_response"] is not None
    assert "analysis_results" in result
```

Migration from Legacy Agents

See [LANGGRAPH-MIGRATION-PLAN.md](#) for comprehensive migration guide.

Quick Migration:

```
# Old (legacy agent coordinator)
from agents.agent_coordinator import AgentCoordinator
coordinator = AgentCoordinator(api_key=api_key)
result = await coordinator.analyze_script(script)

# New (LangGraph orchestrator)
from agents.langgraph_production import LangGraphProductionOrches-
orchestrator = LangGraphProductionOrchestrator(api_key=api_key)
result = await orchestrator.analyze_script(script_content=script)
```

Resources

Documentation

- [LangGraph Official Docs](#)
- [LangChain Documentation](#)
- [Migration Plan](#)

Code Examples

- Production orchestrator: `src/ai/agents/langgraph_production.py`
- API endpoints: `src/ai/langgraph_endpoints.py`
- Tests: `tests/test_langgraph_*.py`

Support

- GitHub Issues: [Repository Issues](#)
 - Team Chat: `#ai-platform`
 - Documentation: `/docs`
-

Document Version: 1.0 **Last Updated:** 2026-01-07 **Maintainer:** AI Platform Team

Generated 2026-01-16 23:34 UTC